

ML/DL for Everyone with PYTORCH

Lecture 9: Softmax Classifier

Sung Kim <hunkim+ml@gmail.com> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>



Call for Comments

Please feel free to add comments directly on these slides.

Other slides: <http://bit.ly/PyTorchZeroAll>



ML/DL for Everyone with PYTORCH

Lecture 9: Softmax Classifier

Sung Kim <hunkim+ml@gmail.com> HKUST


Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>


Videos: <http://bit.ly/PyTorchVideo>




MNIST: 10 labels




10 labels: 10 outputs



10 labels: 10 outputs




10 outputs



$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \underbrace{\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}}_{w \in \mathbb{R}^{2 \times 1}} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}}_{y \in \mathbb{R}^{N \times 1}}$$

10 outputs




$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \underbrace{\begin{bmatrix} w_1 \\ w_2 \end{bmatrix}}_{w \in \mathbb{R}^{2 \times 1}} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}}_{y \in \mathbb{R}^{N \times 1}}$$


$$\underbrace{\begin{bmatrix} a_1 & b_1 \\ a_2 & b_2 \\ \dots & \dots \\ a_n & b_n \end{bmatrix}}_{x \in \mathbb{R}^{N \times 2}} \begin{bmatrix} ? \end{bmatrix} = y \in \mathbb{R}^{N \times 10}$$

$w \in \mathbb{R}^{2 \times ?}$

Probability




Softmax



Meet Softmax

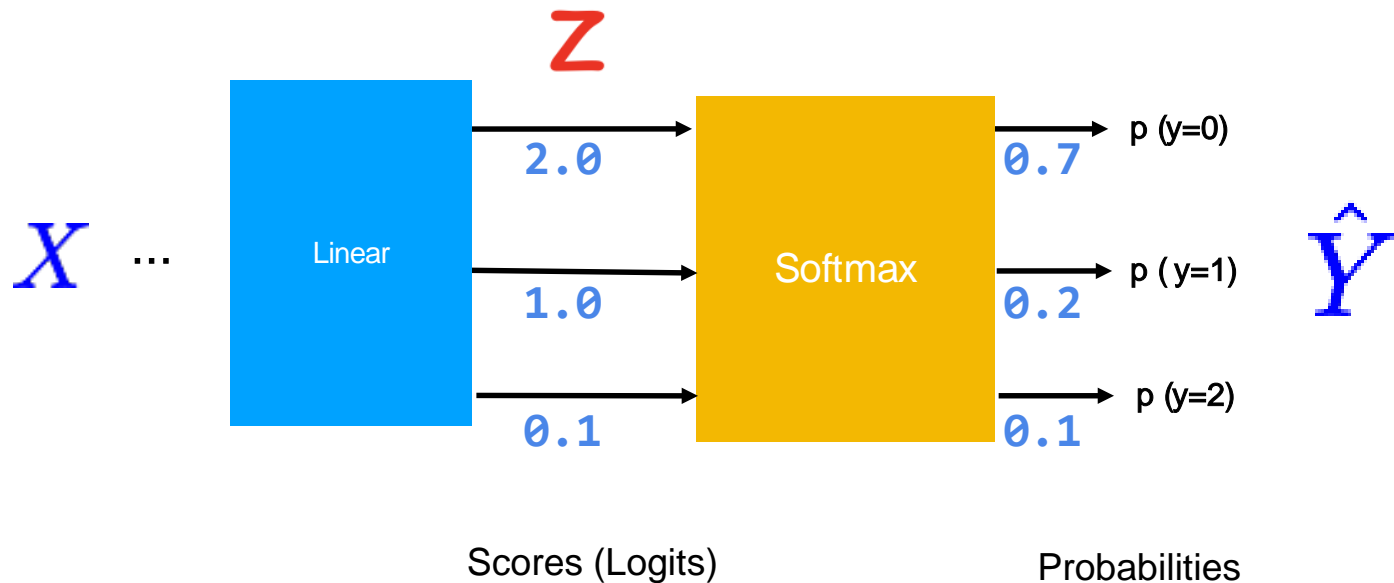
$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$




Scores (Logits)

Meet Softmax

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$





Cost function: cross entropy

A handwritten diagram of the cross-entropy loss formula. The formula is $\mathcal{L} = \frac{1}{N} \sum_i \mathcal{D}(s(w x_i + b), y_i)$. Annotations include: a blue arrow labeled "LOSS" pointing to the \mathcal{L} symbol; a blue arrow labeled "TRAINING SET" pointing to the index i in the summation; and a blue arrow pointing from the y_i term to a large orange \hat{y} symbol above it.


$$\mathcal{L} = \frac{1}{N} \sum_i \mathcal{D}(s(w x_i + b), y_i)$$

$$L(\hat{y}, y) = - \sum_k^K y^{(k)} \log \hat{y}^{(k)}$$

Exercise 9-1: CrossEntropyLoss VS NLLoss


- What are the differences?
- Check out
 - <http://pytorch.org/docs/master/nn.html#nllloss>
 - <http://pytorch.org/docs/master/nn.html#crossentropyloss>
- Minimizing the Negative Log-Likelihood, in English
http://willwolf.io/2017/05/18/minimizing_the_negative_log_likelihood_in_english/

(log)Softmax + NLLLoss



With NLLLoss

MNIST input



$28 \times 28 \text{ pixels} = 784$

Predefined MNIST dataloader

```
batch_size = 64
device = 'cuda' if cuda.is_available() else 'cpu'
print(f'Training MNIST Model on {device}\n{"=" * 44}')

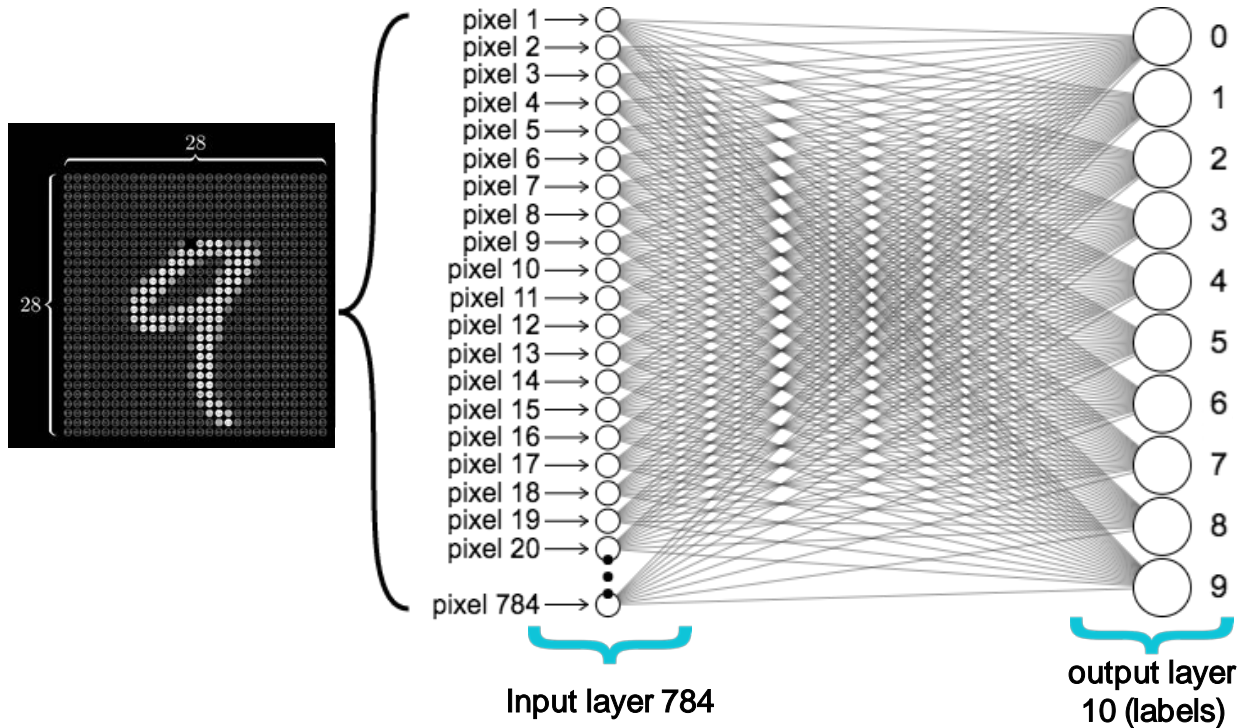
# MNIST Dataset
train_dataset = datasets.MNIST(root='./mnist_data/',
                                train=True,
                                transform=transforms.ToTensor(),
                                download=True)

test_dataset = datasets.MNIST(root='./mnist_data/',
                               train=False,
                               transform=transforms.ToTensor())


# Data Loader (Input Pipeline)
train_loader = data.DataLoader(dataset=train_dataset,
                               batch_size=batch_size,
                               shuffle=True)

test_loader = data.DataLoader(dataset=test_dataset,
                              batch_size=batch_size,
                              shuffle=False)
```

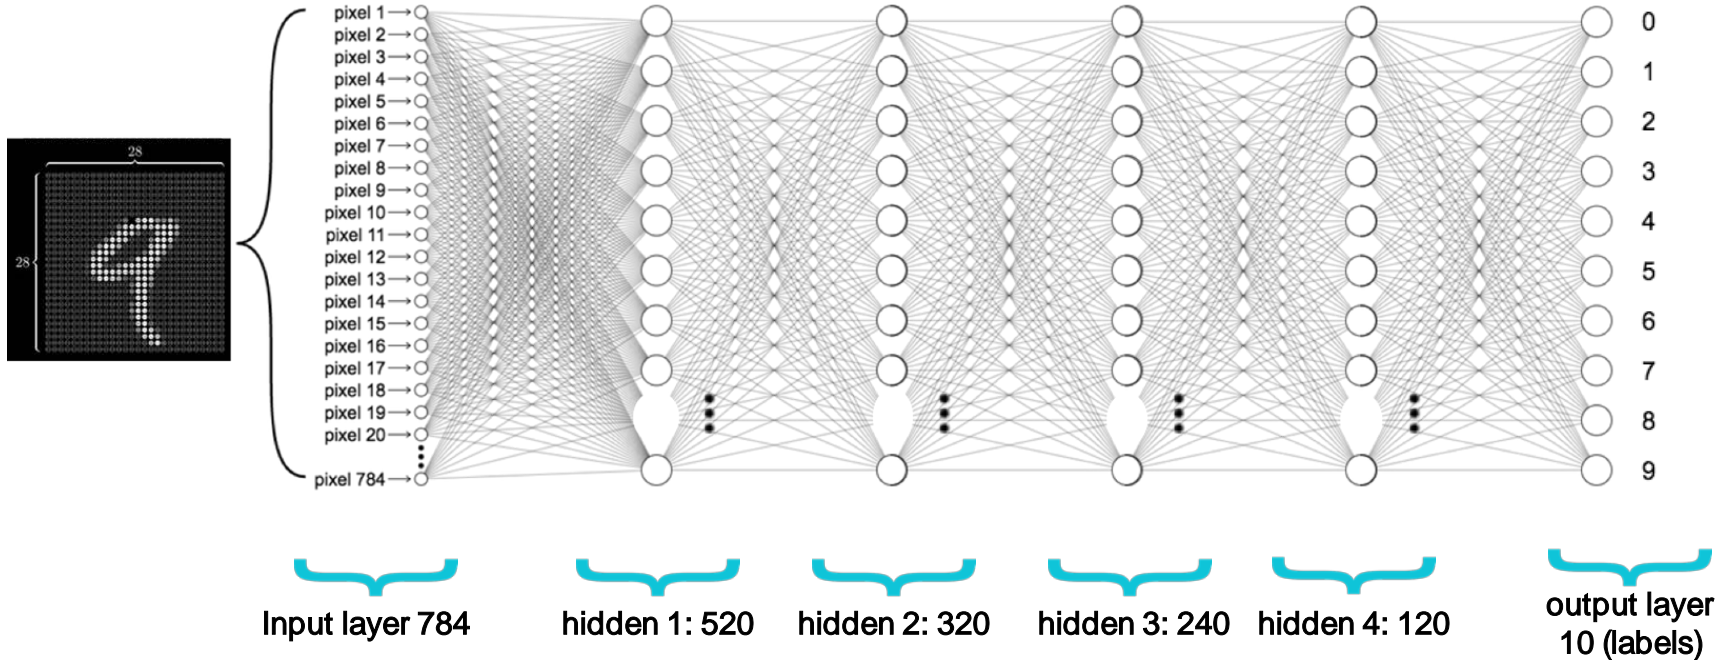
MNIST Network



MNIST Network




MNIST Network





MNIST Network



```
self.l1 = nn.Linear(784, 520)
self.l2 = nn.Linear(520, 320)
self.l3 = nn.Linear(320, 240)
self.l4 = nn.Linear(240, 120)
self.l5 = nn.Linear(120, 10)
```


Softmax & NLL loss



```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        self.l1 = nn.Linear(784, 520)
        self.l2 = nn.Linear(520, 320)
        self.l3 = nn.Linear(320, 240)
        self.l4 = nn.Linear(240, 120)
        self.l5 = nn.Linear(120, 10)
```

```
    def forward(self, x):
        # Flatten the data (n, 1, 28, 28) -> (n, 784)
        x = x.view(-1, 784)
        x = F.relu(self.l1(x))
        x = F.relu(self.l2(x))
        x = F.relu(self.l3(x))
        x = F.relu(self.l4(x))
        return self.l5(x) # No need activation
```



Softmax & NLL loss



```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        self.l1 = nn.Linear(784, 520)
        self.l2 = nn.Linear(520, 320)
        self.l3 = nn.Linear(320, 240)
        self.l4 = nn.Linear(240, 120)
        self.l5 = nn.Linear(120, 10)
```

```
    def forward(self, x):
        # Flatten the data (n, 1, 28, 28)-> (n, 784)
        x = x.view(-1, 784)
        x = F.relu(self.l1(x))
        x = F.relu(self.l2(x))
        x = F.relu(self.l3(x))
        x = F.relu(self.l4(x))
        return self.l5(x) # No need activation
```

```
model = Net()
model.to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)
```


MNIST Train

```
def train(epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        if batch_idx % 10 == 0:
            print('Train Epoch: {} | Batch Status: {}/{} ({:.0f}%) | Loss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

```
def test():
    model.eval()
    test_loss = 0
    correct = 0
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        output = model(data)
        # sum up batch loss
        test_loss += criterion(output, target).item()
        # get the index of the max
        pred = output.data.max(1, keepdim=True)[1]
        correct += pred.eq(target.data.view_as(pred)).cpu().sum()

    test_loss /= len(test_loader.dataset)
    print(f'=====\\nTest set: Average loss: {test_loss:.4f}, Accuracy: {correct}/{len(test_loader.dataset)} '
          f'({100. * correct / len(test_loader.dataset):.0f}%)')
```



Accuracy?

```
if __name__ == '__main__':
    since = time.time()
    for epoch in range(1, 10):
        epoch_start = time.time()
        train(epoch)
        m, s = divmod(time.time() - epoch_start, 60)
        print(f'Training time: {m:.0f}m {s:.0f}s')
        test()
        m, s = divmod(time.time() - epoch_start, 60)
        print(f'Testing time: {m:.0f}m {s:.0f}s')

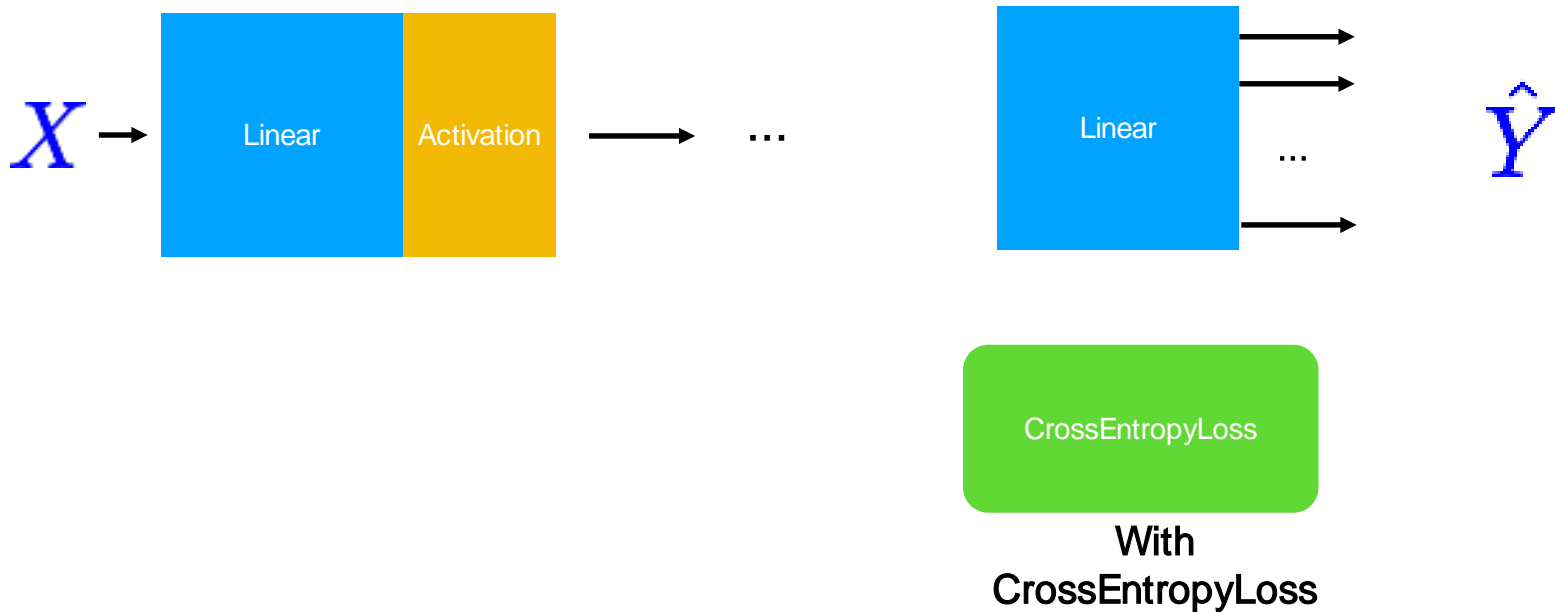
    m, s = divmod(time.time() - since, 60)
    print(f'Total Time: {m:.0f}m {s:.0f}s\nModel was trained on {device}!')
```

Train Epoch: 1 [0/60000 (0%)]	Loss: 2.313209
Train Epoch: 1 [640/60000 (1%)]	Loss: 2.303560
Train Epoch: 1 [1280/60000 (2%)]	Loss: 2.296464
Train Epoch: 1 [1920/60000 (3%)]	Loss: 2.297758
Train Epoch: 1 [2560/60000 (4%)]	Loss: 2.308579
Train Epoch: 1 [3200/60000 (5%)]	Loss: 2.300100
Train Epoch: 1 [3840/60000 (6%)]	Loss: 2.300800
Train Epoch: 1 [4480/60000 (7%)]	Loss: 2.301295
Train Epoch: 1 [5120/60000 (9%)]	Loss: 2.295039
...	
Train Epoch: 9 [51200/60000 (85%)]	Loss: 0.069267
Train Epoch: 9 [51840/60000 (86%)]	Loss: 0.044378
Train Epoch: 9 [52480/60000 (87%)]	Loss: 0.163481
Train Epoch: 9 [53120/60000 (88%)]	Loss: 0.243676
Train Epoch: 9 [53760/60000 (90%)]	Loss: 0.045024
Train Epoch: 9 [54400/60000 (91%)]	Loss: 0.064958
Train Epoch: 9 [55040/60000 (92%)]	Loss: 0.071447
Train Epoch: 9 [55680/60000 (93%)]	Loss: 0.043712
Train Epoch: 9 [56320/60000 (94%)]	Loss: 0.099484
Train Epoch: 9 [56960/60000 (95%)]	Loss: 0.159727
Train Epoch: 9 [57600/60000 (96%)]	Loss: 0.109291
Train Epoch: 9 [58240/60000 (97%)]	Loss: 0.116370
Train Epoch: 9 [58880/60000 (98%)]	Loss: 0.127303
Train Epoch: 9 [59520/60000 (99%)]	Loss: 0.030254

Test set: Average loss: -12.1596, Accuracy: 9697/10000 (97%)

Multiple label prediction?

Just use CrossEntropyLoss!



Exercise 9-2

- Build a classifier for Otto Group Product
 - <https://www.kaggle.com/c/otto-group-product-classification-challenge/data>
 - Use train.csv.zip (1.59 MB)
- Use DataLoader



Lecture 10: CNN