

# ML/DL for Everyone with PYTORCH

## Lecture 11: Advanced CNN

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST

Code: <https://github.com/hunkim/PyTorchZeroToAll>

Slides: <http://bit.ly/PyTorchZeroAll>

Videos: <http://bit.ly/PyTorchVideo>



# Call for Comments

Please feel free to add comments directly on these slides.

Other slides: <http://bit.ly/PyTorchZeroAll>



# ML/DL for Everyone with **PYTORCH**

## Lecture 11: Advanced CNN

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST

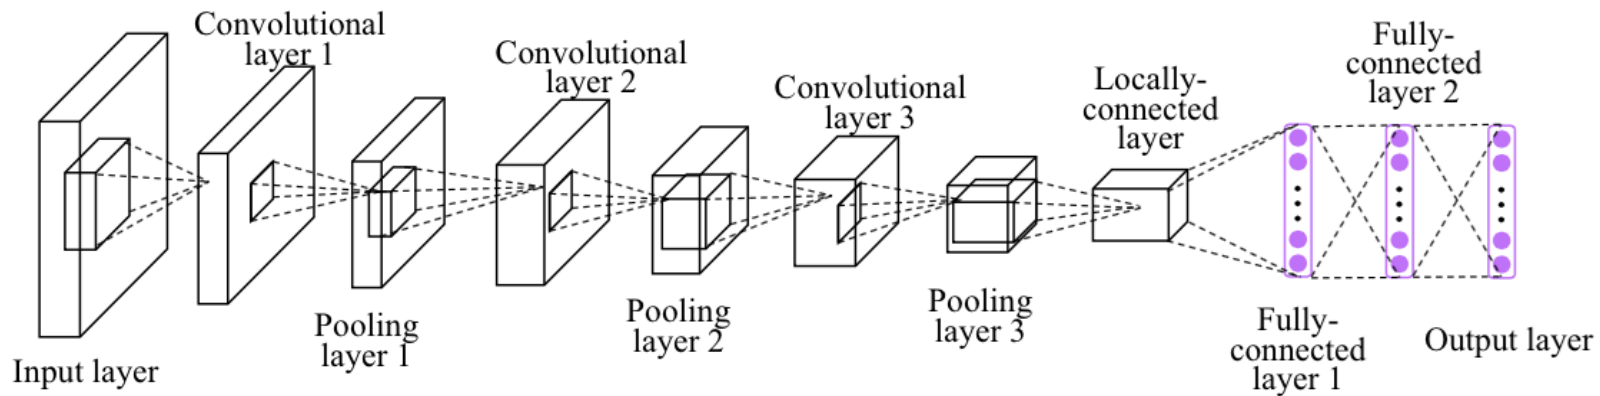
Code: <https://github.com/hunkim/PyTorchZeroToAll>


Slides: <http://bit.ly/PyTorchZeroAll>


Videos: <http://bit.ly/PyTorchVideo>




# CNN








# INCEPTION MODULES

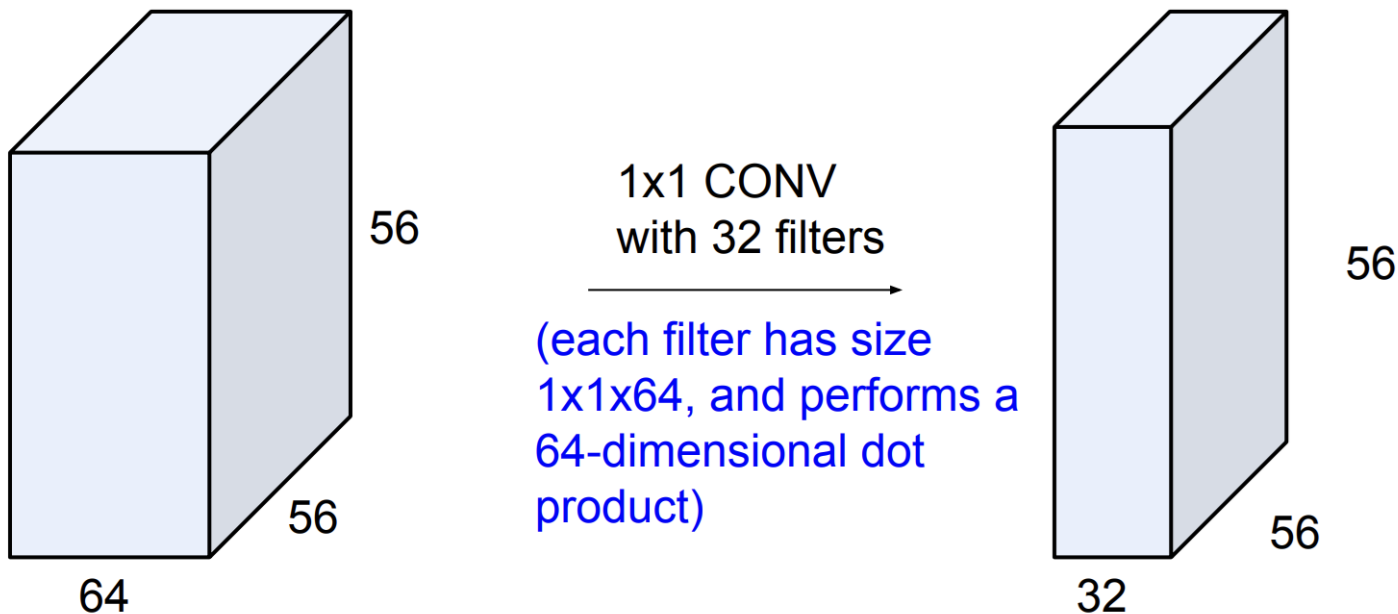


# Why 1x1 convolution?





# Why 1x1 convolution?



# Why 1x1 convolution




Figure 6. 5x5 convolutions inside the Inception module using the naive model




Figure 7. 1x1 convolutions serve as the dimensionality reducers that limit the number of expensive 5x5 convolutions that follow

# Why 1x1 convolution




Figure 6. 5x5 convolutions inside the Inception module using the naive model




Figure 7. 1x1 convolutions serve as the dimensionality reducers that limit the number of expensive 5x5 convolutions that follow

Without 1x1 convolution:  
 $5^2 * 28^2 * 192 * 32 = 120,422,400$  operations

# Why 1x1 convolution




Figure 6. 5x5 convolutions inside the Inception module using the naive model

Without 1x1 convolution:  
 $5^2 * 28^2 * 192 * 32 = 120,422,400$  operations


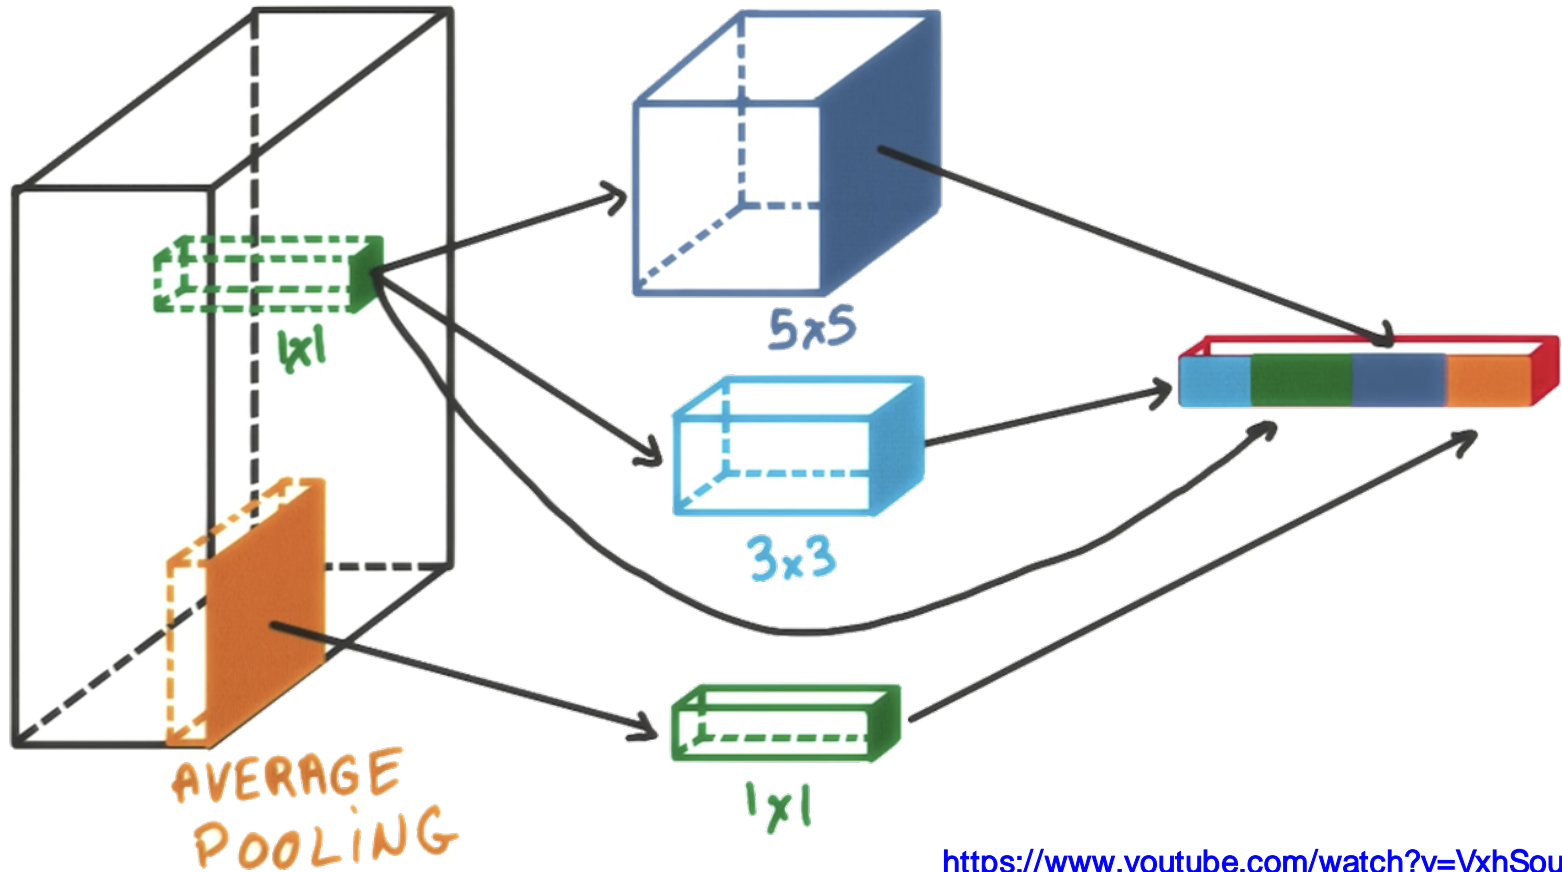


Figure 7. 1x1 convolutions serve as the dimensionality reducers that limit the number of expensive 5x5 convolutions that follow

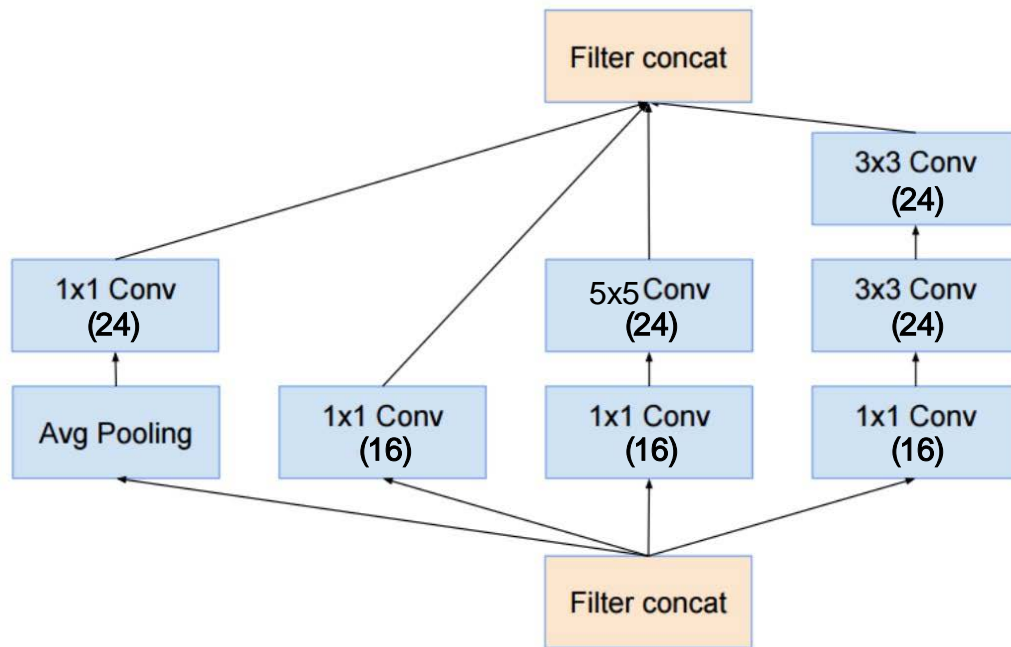
With 1x1 convolution:  
 $1^2 * 28^2 * 192 * 16$   
 $+ 5^2 * 28^2 * 16 * 32$   
 $= 12,443,648$  operations


# INCEPTION MODULES






# Inception Module









```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)  
branch1x1 = self.branch1x1(x)
```





```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)  
  
branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)  
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)  
  
branch1x1 = self.branch1x1(x)
```



```

self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)

```

```

self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

branch1x1 = self.branch1x1(x)


```

```

self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)

branch5x5 = self.branch5x5_1(x)
branch5x5 = self.branch5x5_2(branch5x5)

```



```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

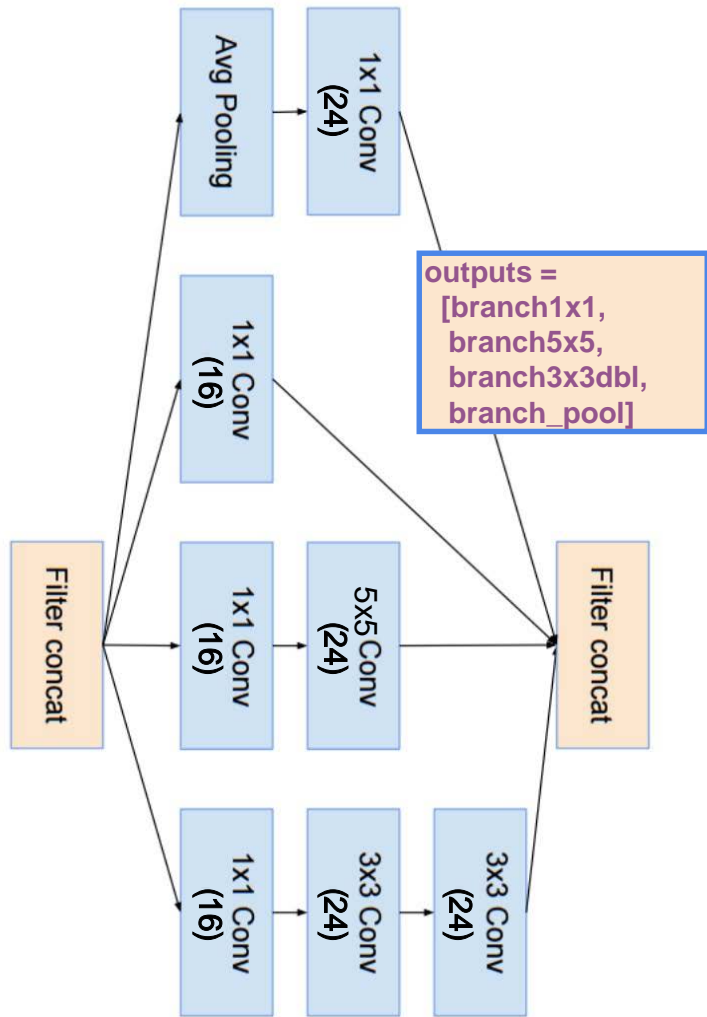
branch1x1 = self.branch1x1(x)
```

```
self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)

branch5x5 = self.branch5x5_1(x)
branch5x5 = self.branch5x5_2(branch5x5)
```

```
self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)

branch3x3dbl = self.branch3x3dbl_1(x)
branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)
```



```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)
```

```
branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
branch1x1 = self.branch1x1(x)
```

```
self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)
```

```
branch5x5 = self.branch5x5_1(x)
```

```
branch5x5 = self.branch5x5_2(branch5x5)
```

```
self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
```

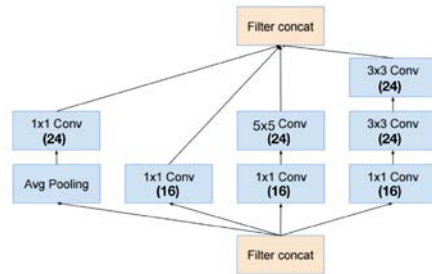
```
self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)
```

```
branch3x3dbl = self.branch3x3dbl_1(x)
```

```
branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
```

```
branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)
```

# Inception Module



```
class InceptionA(nn.Module):
    def __init__(self, in_channels):
        super(InceptionA, self).__init__()
        self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

        self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)

        self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
        self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)

        self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

    def forward(self, x):
        branch1x1 = self.branch1x1(x)

        branch5x5 = self.branch5x5_1(x)
        branch5x5 = self.branch5x5_2(branch5x5)

        branch3x3dbl = self.branch3x3dbl_1(x)
        branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
        branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)

        branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
        branch_pool = self.branch_pool(branch_pool)

        outputs = [branch1x1, branch5x5, branch3x3dbl, branch_pool]
        return torch.cat(outputs, 1)
```

```
class Net(nn.Module):


    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(88, 20, kernel_size=5)

        self.incept1 = InceptionA(in_channels=10)
        self.incept2 = InceptionA(in_channels=20)

        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(1408, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = self.incept1(x)
        x = F.relu(self.mp(self.conv2(x)))
        x = self.incept2(x)
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

# Inception Module



Train Epoch: 9 [44800/60000 (75%)]	Loss: 0.064180
Train Epoch: 9 [45440/60000 (76%)]	Loss: 0.020339
Train Epoch: 9 [46080/60000 (77%)]	Loss: 0.061476
Train Epoch: 9 [46720/60000 (78%)]	Loss: 0.039662
Train Epoch: 9 [47360/60000 (79%)]	Loss: 0.026798
Train Epoch: 9 [48000/60000 (80%)]	Loss: 0.071569
Train Epoch: 9 [48640/60000 (81%)]	Loss: 0.003835
Train Epoch: 9 [49280/60000 (82%)]	Loss: 0.005564
Train Epoch: 9 [49920/60000 (83%)]	Loss: 0.020116
Train Epoch: 9 [50560/60000 (84%)]	Loss: 0.128114
Train Epoch: 9 [51200/60000 (85%)]	Loss: 0.016599
Train Epoch: 9 [51840/60000 (86%)]	Loss: 0.006995
Train Epoch: 9 [52480/60000 (87%)]	Loss: 0.111267
Train Epoch: 9 [53120/60000 (88%)]	Loss: 0.052126
Train Epoch: 9 [53760/60000 (90%)]	Loss: 0.034962
Train Epoch: 9 [54400/60000 (91%)]	Loss: 0.029465
Train Epoch: 9 [55040/60000 (92%)]	Loss: 0.031482
Train Epoch: 9 [55680/60000 (93%)]	Loss: 0.015132
Train Epoch: 9 [56320/60000 (94%)]	Loss: 0.010435
Train Epoch: 9 [56960/60000 (95%)]	Loss: 0.014344
Train Epoch: 9 [57600/60000 (96%)]	Loss: 0.014952
Train Epoch: 9 [58240/60000 (97%)]	Loss: 0.153132
Train Epoch: 9 [58880/60000 (98%)]	Loss: 0.112024
Train Epoch: 9 [59520/60000 (99%)]	Loss: 0.009406

Test set: Average loss: 0.0470, Accuracy: 9866/10000 (**99%**)

```
class Net(nn.Module):

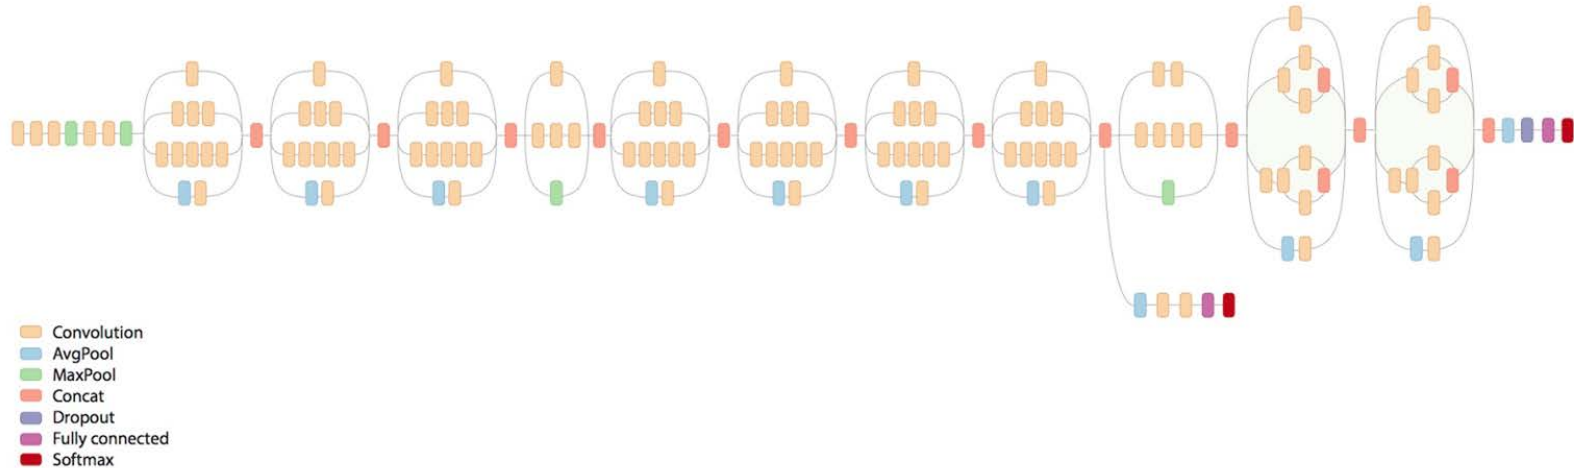
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(88, 20, kernel_size=5)

        self.incept1 = InceptionA(in_channels=10)
        self.incept2 = InceptionA(in_channels=20)


        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(1408, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = self.incept1(x)
        x = F.relu(self.mp(self.conv2(x)))
        x = self.incept2(x)
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

## Exercise 11-1: Implement full inception v3/v4

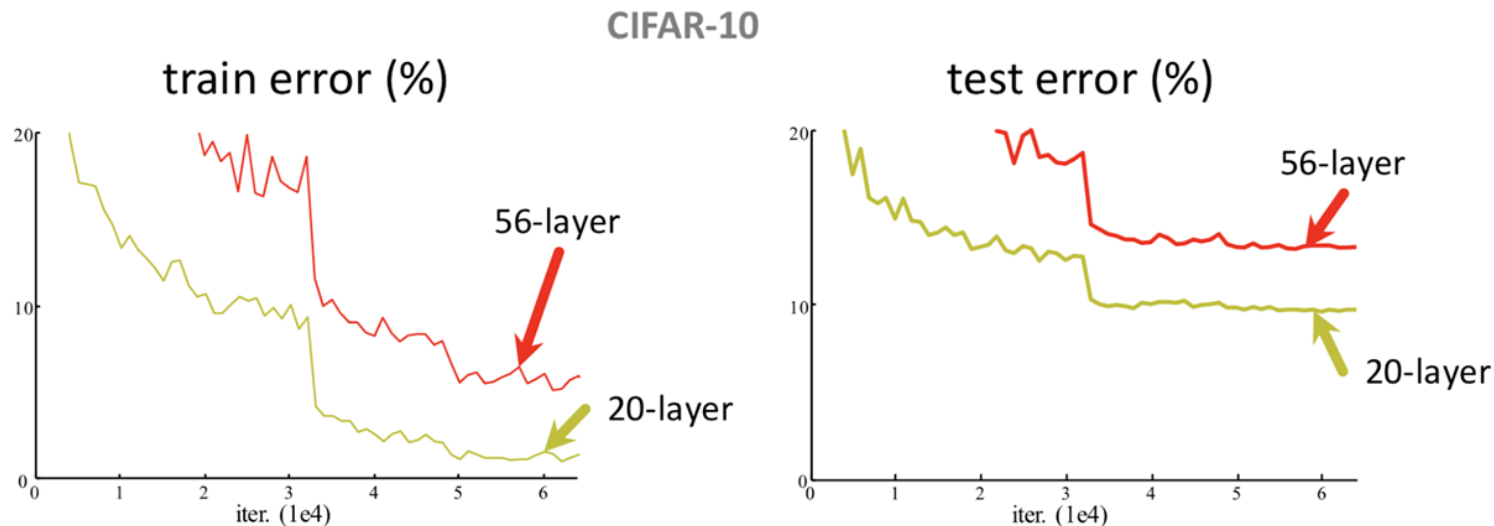


<https://research.googleblog.com/2016/08/improving-inception-and-image.html>







# Can we just go deeper, keep stacking layers?



- *Plain* nets: stacking 3x3 conv layers...
- 56-layer net has **higher training error** and test error than 20-layer net

# Can we just go deeper, keep stacking layers?




- “Overly deep” plain nets have **higher training error**
- A general phenomenon, observed in many datasets

# Problems with stacking layers (TBA)

- Vanishing gradients problem
- Back propagation kind of gives up...
- Degradation problem
  - with increased network depth accuracy gets saturated and then rapidly degrades


# Deep Residual Learning

- Plain net




# Deep Residual Learning

- Plain net




- Residual net




# Network “Design”

- Keep it simple
- Our basic design (VGG-style)
  - all 3x3 conv (almost)
  - spatial size /2 => # filters x2 (~same complexity per layer)
  - **Simple design; just deep!**
- Other remarks:
  - no hidden fc
  - no dropout

plain net




ResNet




# ImageNet experiments

- A practical design of going deeper



**all-3x3**




**bottleneck**  
(for ResNet-50/101/152)


# Many more experiments!

$$h(x) = x$$

error: 6.6%



\* ResNet-110 on CIFAR-10




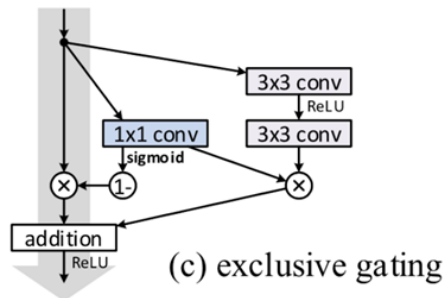
$$h(x) = 0.5x$$

error: 12.4%

$$h(x) = \text{gate} \cdot x$$

error: 8.7%

\*similar to "Highway Network"





$$h(x) = \text{gate} \cdot x$$

error: 12.9%

$$h(x) = \text{conv}(x)$$

error: 12.2%



$$h(x) = \text{dropout}(x)$$

error: > 20%



# ResNets @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**

- ImageNet Classification: “*Ultra-deep*” 152-layer nets
- ImageNet Detection: 16% better than 2nd
- ImageNet Localization: 27% better than 2nd
- COCO Detection: 11% better than 2nd
- COCO Segmentation: 12% better than 2nd

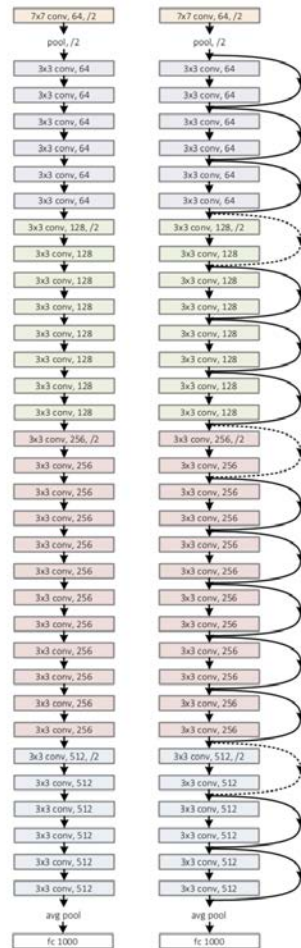
\*improvements are relative numbers

# Exercise 11-2: Implement ResNet


plain net

- Try to implement from scratch
- Deep Residual Learning for Image Recognition:  
<https://arxiv.org/abs/1512.03385>
- Identity Mappings in Deep Residual Networks:  
<https://arxiv.org/abs/1603.05027>
- [http://icml.cc/2016/tutorials/icml2016\\_tutorial\\_deep\\_residual\\_networks\\_kaiminghe.pdf](http://icml.cc/2016/tutorials/icml2016_tutorial_deep_residual_networks_kaiminghe.pdf)

ResNet



# Exercise 11-3: Implement DenseNet



# WHAT NEXT?



## Lecture 12: RNN